



MySQL High Availability at GitHub

Shlomi Noach
GitHub

dataops.barcelona 2018



Agenda

MySQL @ GitHub

The HA story

orchestrator

Old vs. new design

Testing

Thoughts



About me



@github/database-infrastructure

Author of **orchestrator**, **gh-ost**, **freno**, **ccql** and others.

Blog at <http://openark.org>

github.com/shlomi-noach

@ShlomiNoach



GitHub

Built for Developers

Largest open source hosting

67M repositories, 24M users

Critical path in build flows

Best octocat T-Shirts and stickers



MySQL at GitHub

Stores all the metadata: users, repositories, commits, comments, issues, pull requests, ...

Serves web, API and auth traffic

MySQL 5.7, semi-sync replication, RBR, cross DC

~15TB of MySQL tables

~150 production servers, ~15 clusters

Availability is critical



MySQL High Availability



We wish to have:

Automation, reliable detection, DC tolerant failovers, DC isolation tolerance, reasonable failover time, reliable failover, lossless where possible.



MySQL High Availability



Write HA/read HA



MySQL High Availability



Detection

Recovery

Master discovery





orchestrator



orchestrator, meta

Adopted, maintained & supported by
GitHub,


github.com/github/orchestrator

Previously at Outbrain and Booking.com

Orchestrator is free and open source,
released under the Apache 2.0 license

github.com/github/orchestrator/releases

Gaining wider adoption



orchestrator

Discovery


Probe, read instances, build topology graph, attributes, queries

Refactoring


Relocate replicas, manipulate, detach, reorganize

Recovery

Analyze, detect crash scenarios, structure warnings, failovers, promotions, acknowledgements, flap control, downtime, hooks



orchestrator @ GitHub

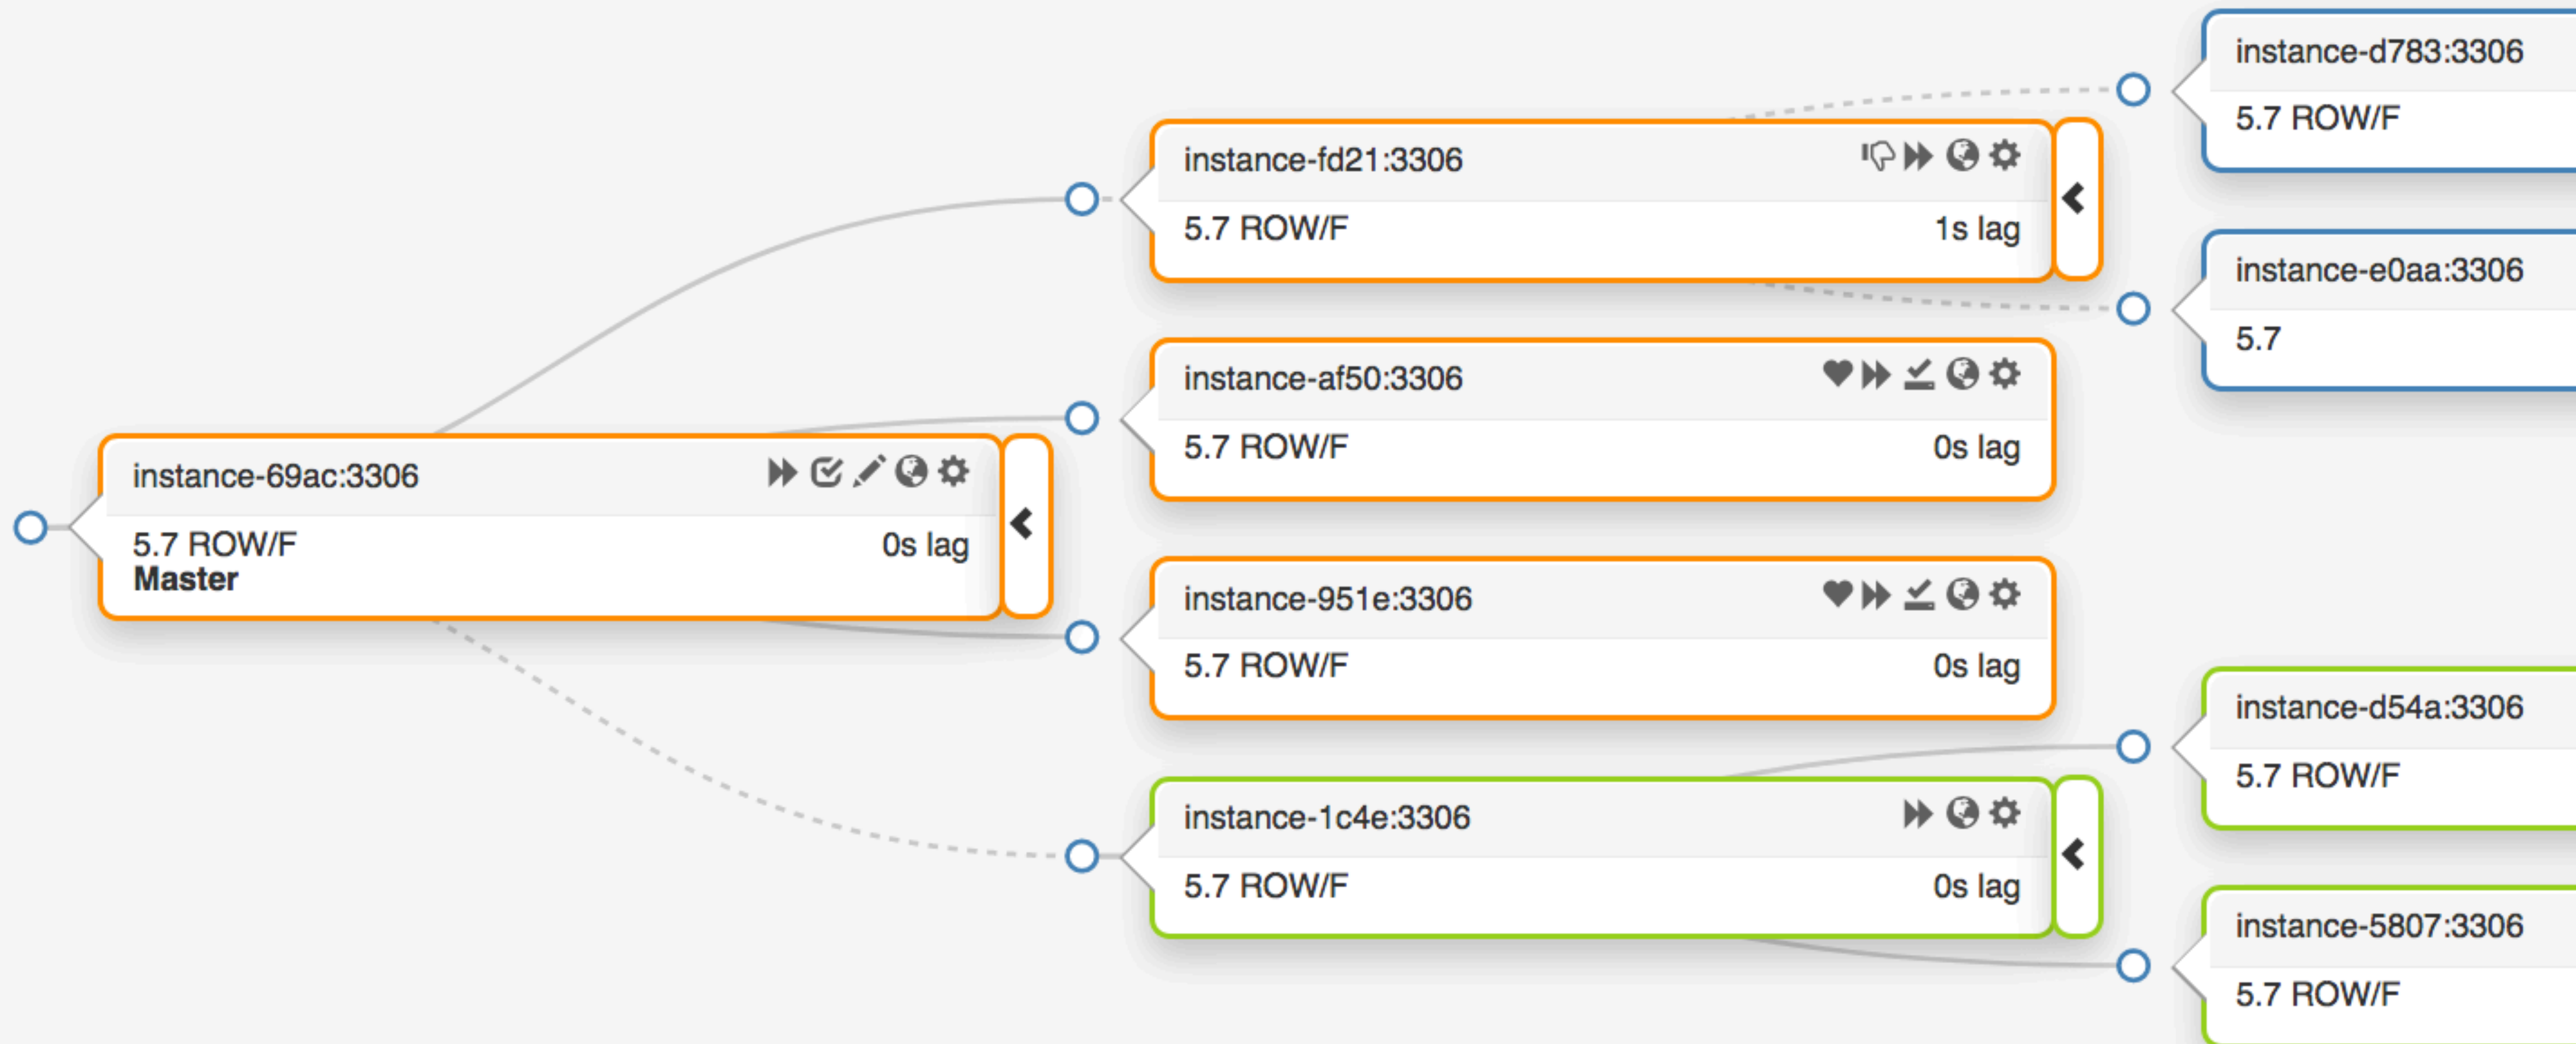


orchestrator/raft deployed on 3 DCs


Automated failover for masters and intermediate masters

Chatops integration





How is orchestrator different?



Holistic approach to failure detection

State based, elaborate recovery decision making




Detection: naive approach

Probe the master

Test failure? Try again n times, interval i



Detection: holistic approach




orchestrator:

Probe the master and its replicas

Expect agreement

Agreement achieved? The cluster is de-facto down.

Promotion: naive assumptions




Configuration indicates which servers whitelisted or blacklisted.

Production operations must reflect in configuration changes.

Promote the most up-to-date replica.



Promotion constraints: real life



orchestrator:

Recognizes environments are dynamic


Understands replication rules

Resolves version, DC, config, promotion rules

Acts based on state



Promotion constraints: real life



orchestrator can promote one, non-ideal replica, have the rest of the replicas converge,
and then *refactor again*, promoting an *ideal* server.



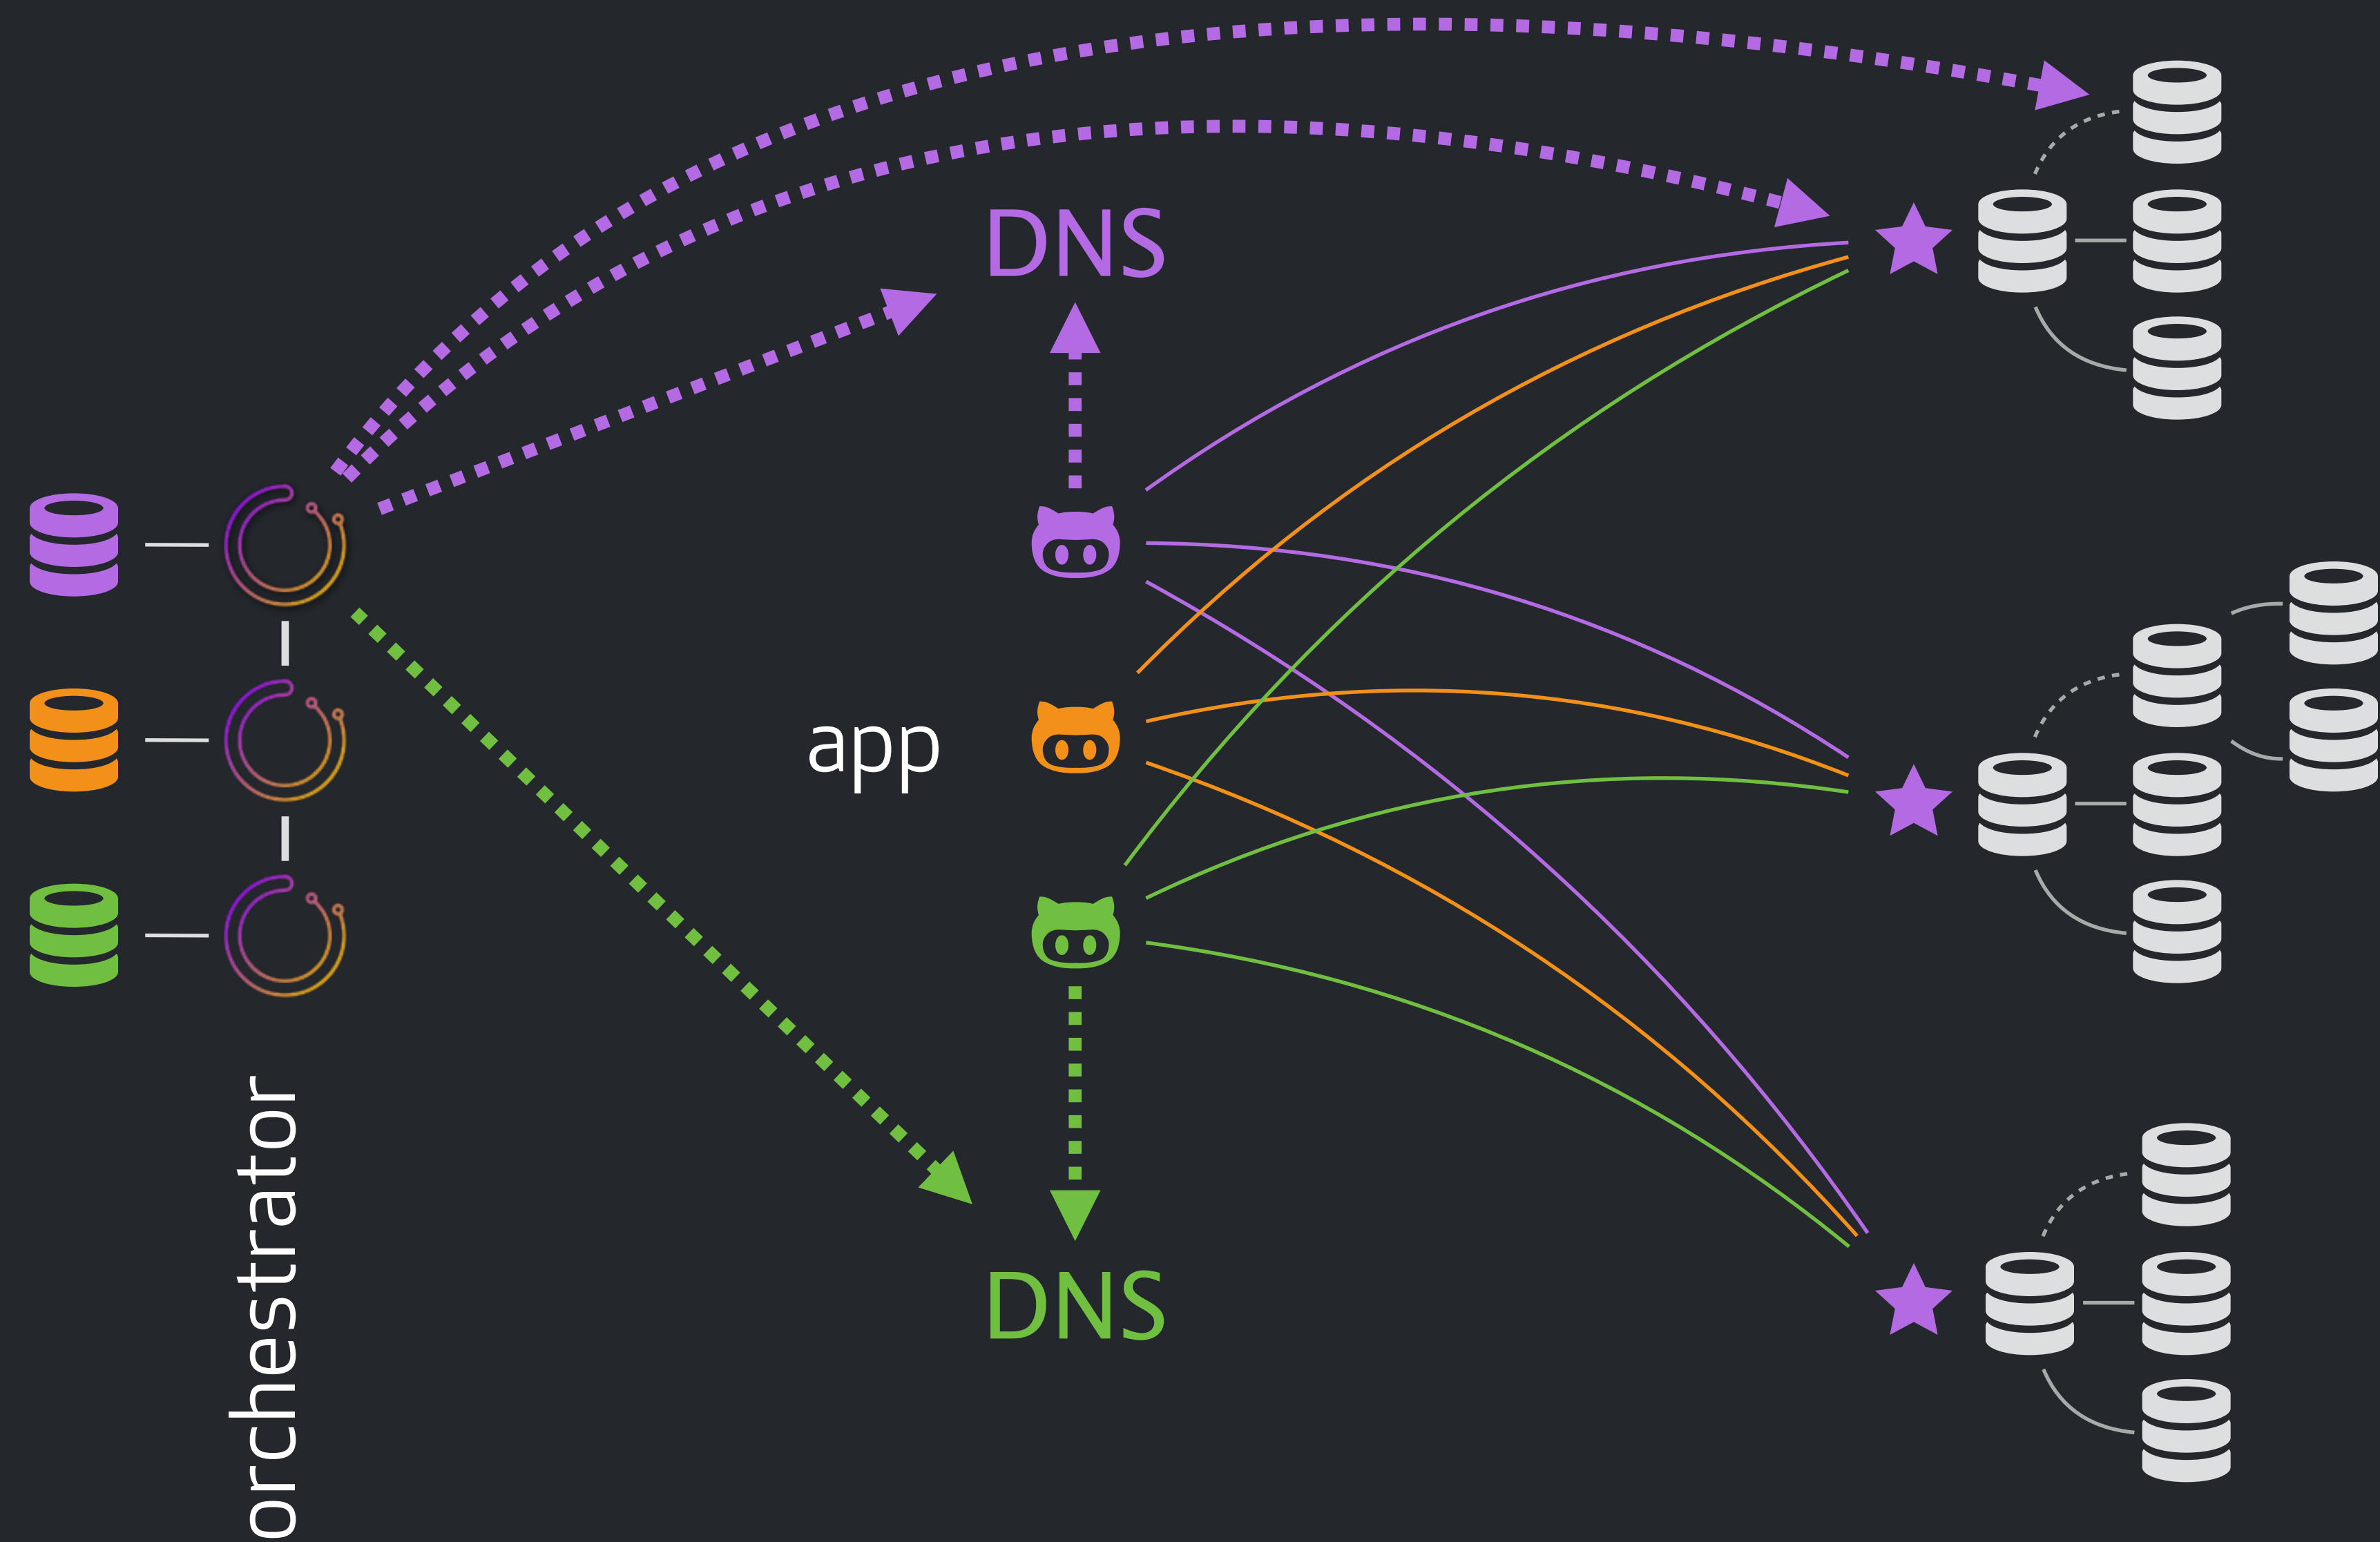
Earlier master discovery @ GitHub



VIP + DNS based



Master discovery via VIP+DNS



Earlier master discovery @ GitHub



Cooperative, long, not a good cross-DC story



A better story



GLB/HAProxy

anycast

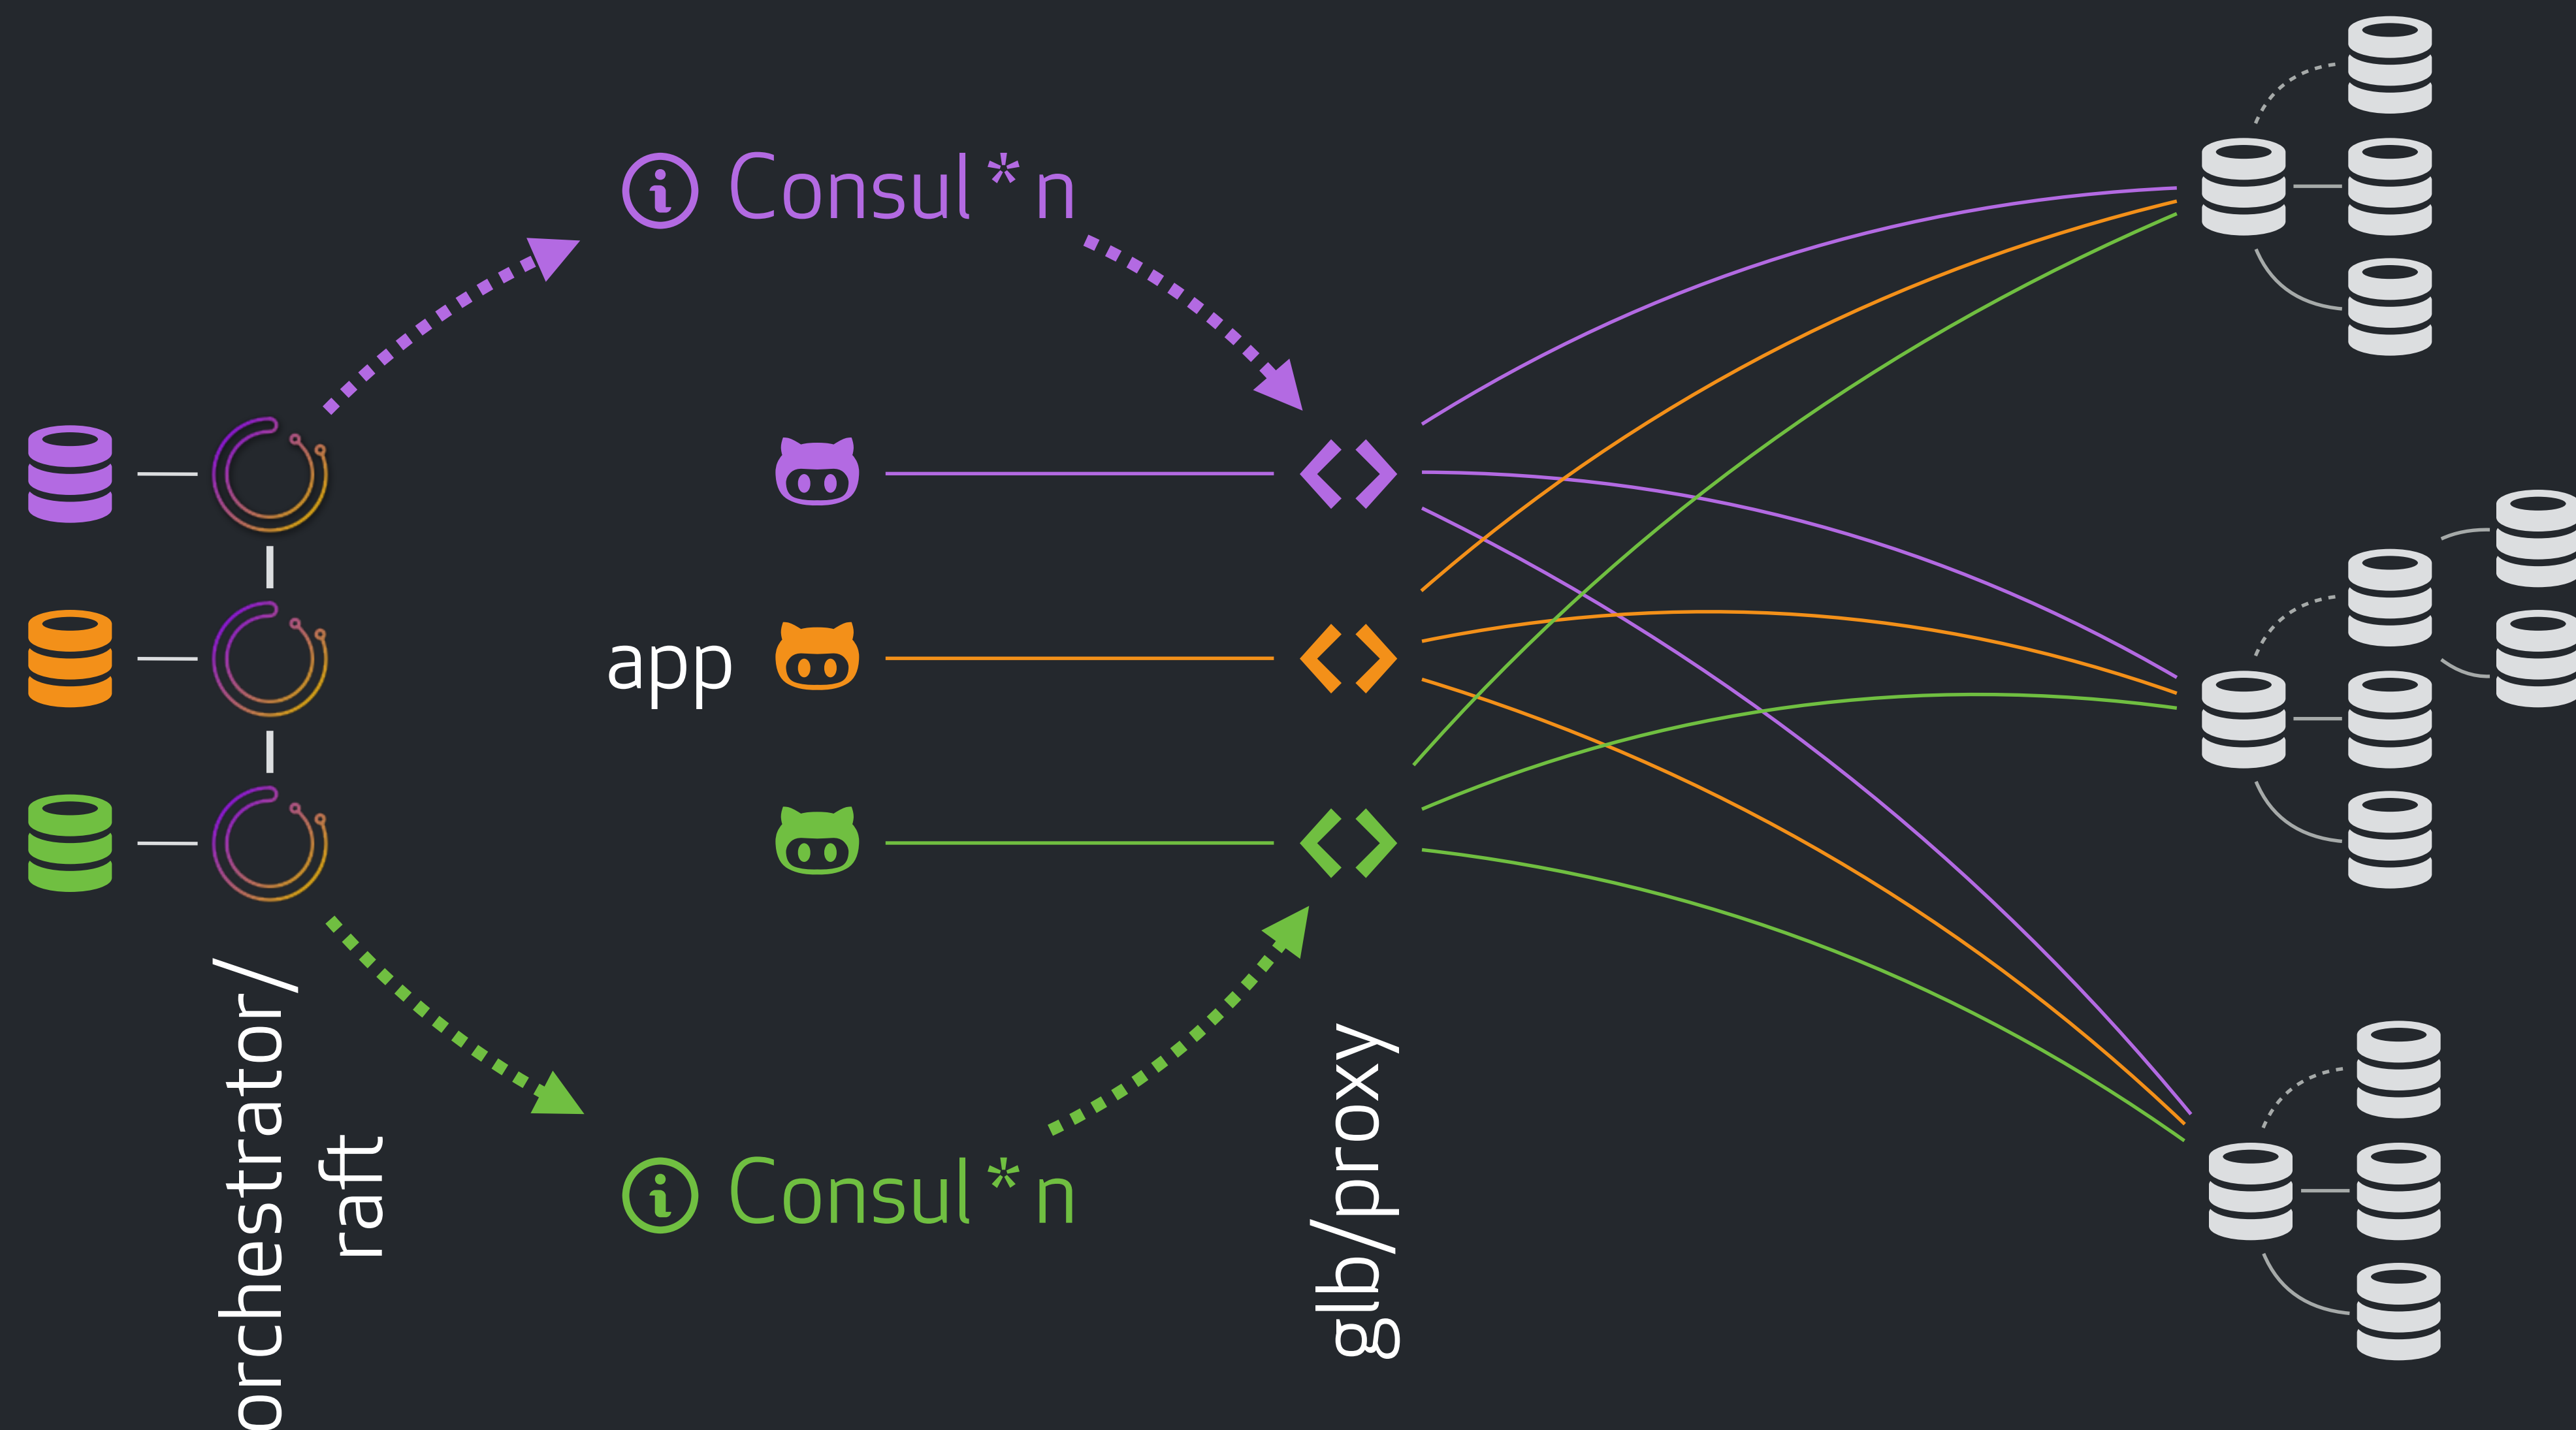
Consul

orchestrator

semi-synchronous replication



orchestrator/Consul/GLB(HAProxy) @ GitHub



A better story



More components, but less moving parts.

Better ownership

Decoupling



➤ with ❤ by **GitHub**



GLB



High available, scalable proxy array

Lossless reloads, implicit SSL, consul integration

GLB director, load balancer array via HAProxy

<https://githubengineering.com/introducing-glb/>

<https://githubengineering.com/glb-part-2-haproxy-zero-downtime-zero-delay-reloads-with-multibinder/>



Consul

By HashiCorp

<https://consul.io/>

Mozilla Public License 2.0

<https://github.com/hashicorp/consul>



Consul

Service Discovery

Health checks, DNS, KV storage

Highly available



consul-template

Simple template engine

Listens to Consul updates




orchestrator/raft

A highly available **orchestrator** setup


Self healing

Cross DC

Mitigates DC partitioning



orchestrator/raft



2 nodes per DC + mediator

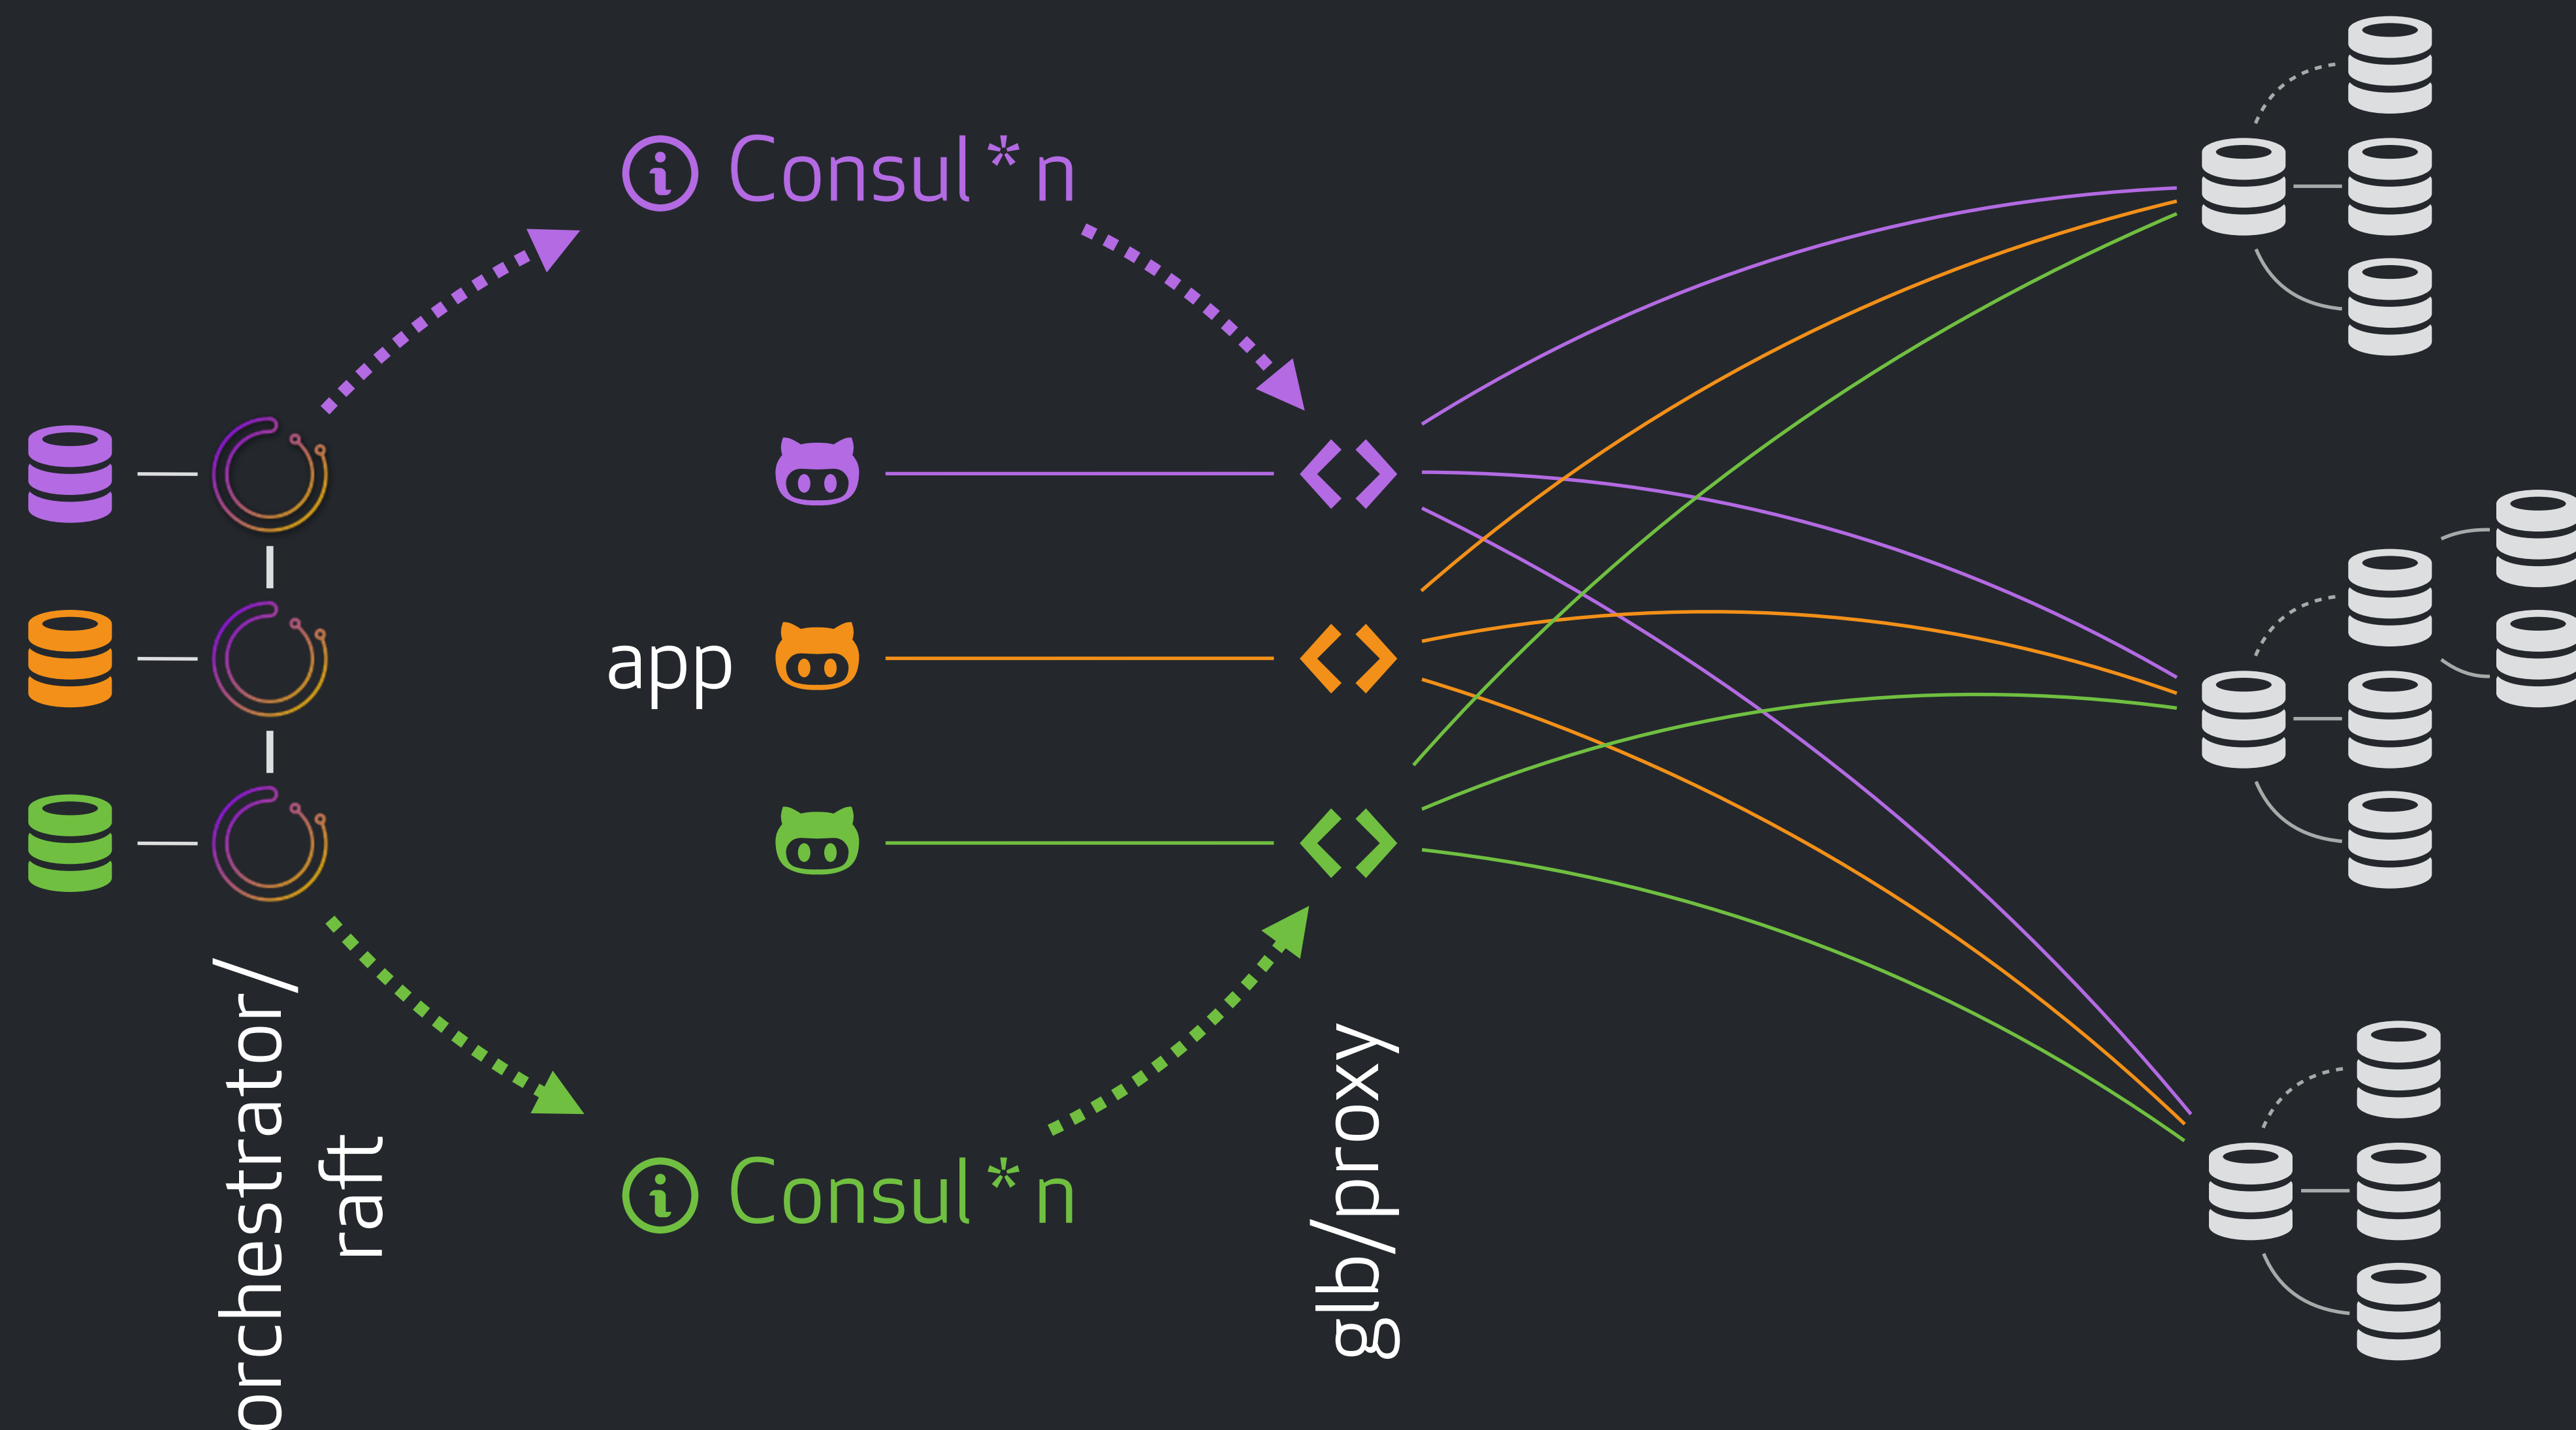
more to come

Added raft features:

- Step down
- Yield
- SQLite log store



orchestrator/Consul/GLB(HAProxy) @ GitHub



orchestrator/Consul/GLB(HAProxy) @ GitHub



orchestrator owns recovery, updates Consul

consul-template runs on GLB servers, reconfigures & reloads GLB

GLB reroutes connections

Hard-kill old connections

Apps connect via anycast, route through local GLB

Independent Consul deployments per DC are managed by orchestrator/raft



pt-heartbeat



Runs on all boxes

Longer poll intervals on **read_only**, supports going in and out of **read_only** mode. Contributed upstream.

<https://github.com/percona/percona-toolkit/pull/302/files?w=1>

No need to start/stop services remotely



Pseudo-GTID



Auto generated by **orchestrator**

Automatically injects on promoted master. No need to start/stop services remotely.



semi-synchronous replication



Lossless, best effort

500ms timeout

Effectively picks our ideal candidates



Results



Reliable detection

Recovery in:

10s - 13s (total outage time), normal case

15s - 20s, difficult case

25s, rare



Cons



App identity unknown

Distributed system, calls for a variety of scenarios

STONITH, work in progress



Testing



Testing cluster in production environment

Continuously kill/block/reject



Thoughts



STONITH

Retries

orchestrator + Consul + proxy as appliance

Kubernetes



Thank you!



Questions?

github.com/shlomi-noach

@ShlomiNoach

